

Lecture 6 — Processes in UNIX

Jeff Zarnett

`jzarnett@uwaterloo.ca`

Department of Electrical and Computer Engineering
University of Waterloo

April 26, 2015

In UNIX a process may create other processes.

The creating process is the parent; newly-created is the child.

Every process has a parent, stretching back to `init`.

Each process has a unique identifier in its process control block.

This is the `pid` (process ID).

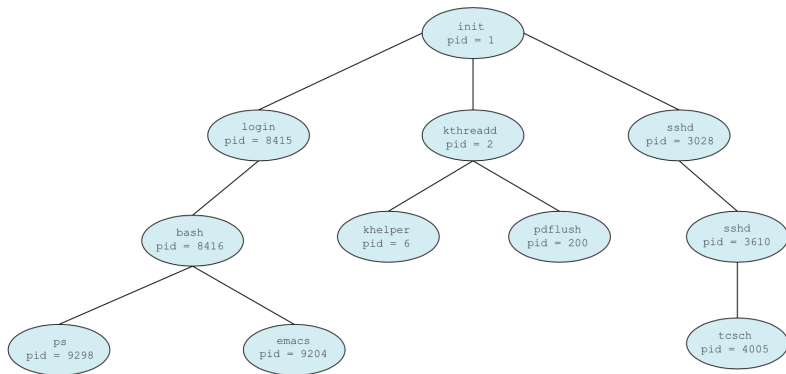
For the most part, users will not need to know or think about the ID.

Exception when trying to terminate one that's gotten stuck.
(`kill -9 24601`).

The `init` process always gets a `pid` of 1.

I don't recommend trying to kill `init`.

Linux Process Tree



We can obtain a list of processes with `ps`.

The diagram shows each user gets a `login` process.

The shell (`bash`) is spawned from `login`.

When you issue a command, like `ls` or `top` (table of processes), the new process is created and the shell will wait on that process.

It might finish on its own (e.g., `ls`).

Or wait for the user to tell it to exit (`top`)

When it does, control goes back to the shell.

You get presented with the prompt again (e.g., `jz@Loki: ~/$`).

Must I log in to the system in a second terminal window to run two things at a time?

The answer is no, and there are two ways to get around it.

Option 1: tell the shell we want the task to run in the background.

To do that, add to the command the `&` symbol:

```
gcc fork.c &
```

Control returns almost immediately to the shell.
It is not waiting for `gcc` to finish.

You may see some output like `[1] 34429`.

This is the shell saying: child has been created; it has process ID 34429.

When the process is finished, there is another update:

```
[1]+ Done gcc fork.c
```


Notably, any console output that the `gcc` command would generate will still appear on the console where the background task was created.

Maybe you want that but maybe you want to put the output in a log file, with a command like `cat fork.c > logfile.txt &`.

(Telling `gcc` to be silent is a somewhat more complex operation.)

The semantics of `&` are not just “run this in the background, please”.

It is actually the parent process (the shell) disowning its child.

That process will get adopted by `init`.

It can run to completion even if the user logs out.

A common example of a command I use involving the &:

```
sudo service xyz start &
```

This will (with super user permissions) start up the service xyz.

It returns control to the console so I don't have to wait.

Next: `tail -f /var/log/xyz/console.log`

Watch the console log of the xyz service as it starts up.

The other alternative is the `screen` command.

While having something run in the background is nice, it does not work for interactive processes.

Example: text editing with `vi` and want to read e-mail with `pine`.

Could be done by saving and closing `vi`.

Or, start them in `screen` and switch between them.

Instead of just opening `vi fork.c` I can issue the command `screen vi fork.c` and this spawns `screen` and takes me right to editing the file.

The key difference is that I can “detach” from this screen and go back to the command line.

If I log out, `screen` keeps running with the `vi` inside it.

If I have multiple screens running, I can just “reattach” to the one I want.

In general, when a process spawns a child, the child will need resources.

The child may request them from the OS directly.

Or the parent can give some of its resources to the child.

The parent may partition resources amongst the children or allow its children to share.

Restrict a child process to only some subset of its parent's resources?

If so, cannot overload the system by spawning too many children.

At the time of initialization, the parent may pass the child some data.

Example: link from e-mail to browser.

Interesting note: child may be a duplicate or totally new.

Parent spawns the child process with the `fork` system call.

If waiting for the child process to finish, `wait`.
Alternatively, carry on.

When the child process is finished, it returns a value with `exit`

The parent gets this as the return value of `wait` and may proceed.

Note: `fork` creates a new process as a copy of itself.

Both parent and child continue after that statement.

The call `fork` can return a value:

- A negative value means the fork failed.

- A zero value means this process is the child.

- A positive value: this is the parent; the value is the child `pid`.

After the `fork`, one of the processes may use the `exec` system call.

This will replace its memory space with a new program.

There's no rule that says this must happen
a child can continue to be a clone of its parent if it wishes.

The `exec` invocation loads a binary file into memory & starts execution.

At this point, the programs can go their separate ways.

Or the parent might want to wait for the child to finish.

```
int main()
{
    pid_t pid;
    int childStatus;

    /* fork a child process */
    pid = fork();

    if (pid < 0) {

        /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;

    } else if (pid == 0) {
        /* child process */
        execlp("/bin/ls", "ls", NULL);


    } else {
        /* parent process */
        /* parent will wait for the child to complete */
        wait(&childStatus);
        printf("Child Complete with status: %i \n", childStatus);

    }
    return 0;
}
```

Thus, the output is:

```
jz@Freyja:~/fork$ ./fork
fork    fork.c
Child Complete with status: 0
jz@Freyja:~/fork$
```

Or, to represent this visually:



What about termination?

On the assumption that the process is terminating normally and not being killed, the system call for that is `exit`.

If the program itself has no explicit call to `exit`, the `return` statement at the end of `main` will have the same effect.

Let us modify that code above to fork off a child process that will exit “abnormally” with an exit code of 1.

The `wait` function also returns the process ID of the child.

This is so that the parent can identify which of its children has terminated, though it is not used in this example.

```
int main()
{
    pid_t pid;
    int childStatus;

    /* fork a child process */
    pid = fork();

    if (pid < 0) {

        /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;

    } else if (pid == 0) {
        /* child process */
        exit( 1 );

    } else {
        /* parent process */
        /* parent will wait for the child to complete */
        wait(&childStatus);
        printf("Child Complete with status: %i \n", childStatus);

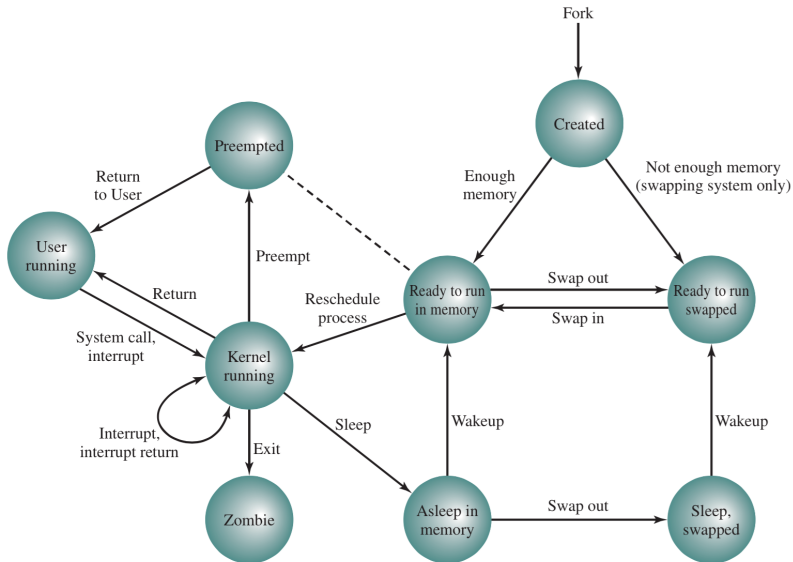
    }
    return 0;
}
```

UNIX divides its processes into two categories: system processes that run in kernel mode and user processes that run in user mode.

There are nine different states:

- 1 User Running**
- 2 Kernel Running**
- 3 Ready to Run, in Memory**
- 4 Asleep in Memory**
- 5 Ready to Run, Swapped**
- 6 Sleeping, Swapped**
- 7 Preempted**
- 8 Created**
- 9 Zombie**

UNIX Process States



Process creation when `fork` is called means the OS does the following:

- 1 It allocates a slot in the process table for the new process.
- 2 It assigns a unique process ID to the child process.
- 3 It makes a copy of the process image of the parent, with the exception of any shared memory.
- 4 It increments counters for any files owned by the parent (showing there is an additional process referencing those files).
- 5 The new process is in the state Ready to Run.
- 6 A return value of 0 goes to the child process, and the unique process ID of the child is returned to the parent.

Afterwards, the system will need to choose which process is going to run:

- 1 The parent process. The child is in the ready to run state.
- 2 The child process. The parent is in the ready to run state.
- 3 Another process. Both parent and child are in the ready to run state.

A short digression on a denial of service attack: the “fork bomb”.

The idea is to call `fork` repeatedly.

Keep doing this until the system crashes (or no work can get done).

Exponential growth (2^n) processes after n calls.

A system can be configured to defend against this.

1. Limit total number of processes per user.
2. Limit rate of process spawning.

Note: do not attempt this on University computers!